87.5%

Congratulations, you passed!

Great job on the challenge, you just earned rewards to use on more challenges.

Add to profile (https://www.linkedin.com/profile/add?startTask=QUIZ:BIGQUERY)

Linux Academy

Back to Lesson

Report Card

Expectations	Score
1. google cloud data engineer - bigquery	87.5 %

Exam Breakdown

google cloud data engineer - bigquery

1. **Which Hadoop ecosystem service is most direct equivalent to BigQuery?**

A. HDFS	
B. PySpark	
C. Spark	
D. Hive	✓ **Correct**

2. **What are the different partitioning methods on BigQuery? (Choose two)**

A. Ingestion-time partitioned	✓ **Correct**
B. Conditional partitions	
C. Sharding date	
D. Partition by specific timestamp/date column	✓ **Correct**

3. **You have a very large table with many columns that are not immediately relevant to your non-IT team members. You want to reduce the amount of irrelevant column data available in your table in order to keep from confusing your team members that need to run queries against it. What is a valid method of achieving this task?**

A. Create a copy of the table, and delete the unneeded columns.	
B. Assign IAM roles to only the columns that your staff needs access to.	
C. Use views to restrict the amount of data available, and have team members run their query against the view.	✓ **Correct**
D. Hide the unneeded columns from the table.	

4. **You are viewing the details of a recent large query and notice that Stage 1 has a full purple bar. What does this tell you?**

A. Stage 1 spent most of its time reading from a large dataset. **Why is this correct?** The purple indicator indicates high read time. https://linuxacademy.com/cp/courses/lesson/course/2238/lesson/4/module/208 (https://linuxacademy.com/cp/courses/lesson/course/2238/lesson/4/module/208)	✓ **Correct**
B. Stage 1 had a heavy computation procedure	
C. Your query could not be successfully completed.	
D. Your query stage spent a long time waiting for additional input. **Why is this incorrect?** This is not correct. https://linuxacademy.com/cp/courses/lesson/course/2238/lesson/4/module/208 (https://linuxacademy.com/cp/courses/lesson/course/2238/lesson/4/module/208)	✕ **Your Answer**

5. **Choose two best practices for creating more efficient queries and saving costs.**

A. Save your biggest JOINs for last	
B. Avoid using `SELECT *` for column selection.	✓ **Correct**
C. Normalize your data when possible.	
D. Filter early and big with `WHERE` clauses	✓ **Correct**

6. **In order to save on bandwidth costs, you need to load data into BigQuery in a compressed state. What is the preferred file format to balance size and performance?**

A. Avro	✓ **Correct**
B. PY	
C. JSON	
D. CSV	

7. Which of these is *not* a valid BigQuery data format?

A. JSON	
B. AVRO	
C. CSV	
D. DOC	✓ **Correct**

8. Which of these actions can you *not* perform with the BigQuery Web UI?

A. Assign IAM roles.	
B. Toggle Legacy SQL to *off* on queries.	
C. Load multiple files at once	✓ **Correct**
D. Create dataset	